

# Data Repository Service

# Table of Contents

1. Overview .....	1
1.1. Version information .....	1
1.2. Contact information .....	1
1.3. License information .....	1
1.4. URI scheme .....	1
1.5. Tags .....	1
1.6. Consumes .....	1
1.7. Produces .....	1
2. Introduction .....	2
3. DRS API Principles .....	3
3.1. DRS IDs .....	3
3.2. DRS Datatypes .....	3
3.3. Read-only .....	3
3.4. URI convention .....	3
3.5. Standards .....	4
4. Authorization & Authentication .....	5
4.1. Making DRS Requests .....	5
4.2. Fetching DRS Objects .....	5
5. Motivation .....	6
5.1. Federation .....	7
6. Paths .....	8
6.1. Get info about a Data Bundle. ....	8
6.1.1. Description .....	8
6.1.2. Parameters .....	8
6.1.3. Responses .....	8
6.1.4. Tags .....	8
6.2. Get info about a Data Object. ....	8
6.2.1. Description .....	8
6.2.2. Parameters .....	8
6.2.3. Responses .....	9
6.2.4. Tags .....	9
6.3. Get a URL for fetching bytes. ....	9
6.3.1. Description .....	9
6.3.2. Parameters .....	9
6.3.3. Responses .....	9
6.3.4. Tags .....	10
6.4. Get information about this implementation. ....	10
6.4.1. Description .....	10

6.4.2. Responses .....	10
6.4.3. Tags .....	10
7. Definitions .....	11
7.1. AccessMethod .....	11
7.2. AccessURL .....	11
7.3. Bundle .....	11
7.4. BundleObject .....	13
7.5. Checksum .....	13
7.6. Error .....	13
7.7. Object .....	13
7.8. ServiceInfo .....	14

# Chapter 1. Overview

<https://github.com/ga4gh/data-repository-service-schemas>

## 1.1. Version information

*Version* : 0.1.0

## 1.2. Contact information

*Contact* : GA4GH Cloud Work Stream

*Contact Email* : [ga4gh-cloud@ga4gh.org](mailto:ga4gh-cloud@ga4gh.org)

## 1.3. License information

*License* : Apache 2.0

*License URL* : <https://raw.githubusercontent.com/ga4gh/data-repository-service-schemas/master/LICENSE>

*Terms of service* : <https://www.ga4gh.org/terms-and-conditions/>

## 1.4. URI scheme

*BasePath* : /ga4gh/drs/v1

*Schemes* : HTTPS

## 1.5. Tags

- DataRepositoryService

## 1.6. Consumes

- `application/json`

## 1.7. Produces

- `application/json`

# Chapter 2. Introduction

The Data Repository Service (DRS) API provides a generic interface to data repositories so data consumers, including workflow systems, can access data in a single, standard way regardless of where it's stored and how it's managed. This document describes the DRS API and provides details on the specific endpoints, request formats, and responses. It is intended for developers of DRS-compatible services and of clients that will call these DRS services.

The primary functionality of DRS is to map a logical ID to a means for physically retrieving the data represented by the ID. The sections below describe the characteristics of those IDs, the types of data supported, and how the mapping works.

# Chapter 3. DRS API Principles

## 3.1. DRS IDs

Each implementation of DRS v1 can choose its own id scheme, as long as it follows these guidelines:

- DRS IDs are URL-safe text strings made up of alphanumeric characters and any of [`._/`]
- One DRS ID MUST always return the same object data (or, in the case of a collection, the same set of objects). This constraint aids with reproducibility.
- DRS does NOT support semantics around multiple versions of an object. (For example, there's no notion of “get latest version” or “list all versions” in DRS v1.) Individual implementation MAY choose an ID scheme that includes version hints.
- DRS implementations MAY have more than one ID that maps to the same object.

## 3.2. DRS Datatypes

DRS v1 supports two types of content:

- an **Object** is like a file — it's a single blob of bytes
- a **Bundle** is like a folder — it's a collection of other DRS content (either objects or bundles)

## 3.3. Read-only

DRS v1 is a read-only API. We expect that each implementation will define its own mechanisms and interfaces (graphical and/or programmatic) for adding and updating data.

## 3.4. URI convention

For convenience, we define a recommended syntax for fully referencing DRS-accessible objects. Strings of the form `drs://<server>/<path>/<[bundle | object]>/<endpoint>/<id>` mean “make a DRS GET to the HTTPS address at `<server>`, passing in the DRS id `<id>` at the end of the appropriate `<path>/<[bundle | object]>/<endpoint>`, to retrieve the DRS JSON response”.

The `<endpoint>` is optional depending on what you are doing e.g. the `<endpoint>` is `/access` when retrieving a particular access method. This is all just

a complicated way of just saying “replace `drs://` with `https://` and GET on the URL that results.”

For example:

- `drs://test.drs.org/ga4gh/drs/v1/bundles/5007fc21-6430-4fb9-82b3-520c60e26751` → maps to a GET request to <https://test.drs.org/ga4gh/drs/v1/bundles/5007fc21-6430-4fb9-82b3-520c60e26751>
- `drs://test.drs.org/ga4gh/drs/v1/objects/65a3e501-7448-4887-aada-b52de5c8db31` → maps to a GET request to <https://test.drs.org/ga4gh/drs/v1/objects/65a3e501-7448-4887-aada-b52de5c8db31>
- `drs://test.drs.org/ga4gh/drs/v1/objects/65a3e501-7448-4887-aada-b52de5c8db31/access/1` → maps to a GET request to <https://test.drs.org/ga4gh/drs/v1/objects/65a3e501-7448-4887-aada-b52de5c8db31/access/1>

[b52de5c8db31/access/1](#)

The service-info endpoint can be queried to retrieve information about the version of DRS supported by this server e.g. <https://<server>/<path>/service-info>

These URIs will be useful when passing objects to a WES server for processing since they provide a way to understand how to successfully make an HTTPS request and process the JSON response.

## 3.5. Standards

The DRS API specification is written in OpenAPI and embodies a RESTful service philosophy. It uses JSON in requests and responses and standard HTTPS for information transport.

# Chapter 4. Authorization & Authentication

## 4.1. Making DRS Requests

The DRS implementation is responsible for defining and enforcing an authorization policy that determines which users are allowed to make which requests. We recommend that DRS implementations use an OAuth2 [bearer token](#), although they can choose other mechanisms if appropriate. The [service-info](#) endpoint should provide sufficient information for a user to figure out how to authenticate with a DRS implementation.

## 4.2. Fetching DRS Objects


The DRS API allows implementers to support a variety of different content access policies, depending on what [AccessMethod](#)s they return:

- public content:
  - server provides an [access\\_url](#) with a [url](#) and no [headers](#)
  - caller fetches the object bytes without providing any auth info
- private content that requires the caller to have out-of-band auth knowledge (e.g. service account credentials):
  - server provides an [access\\_url](#) with a [url](#) and no [headers](#)
  - caller fetches the object bytes, passing the auth info they obtained out-of-band
- private content that requires the caller to pass an Authorization token:
  - server provides an [access\\_url](#) with a [url](#) and [headers](#)
  - caller fetches the object bytes, passing auth info via the specified header(s)
- private content that uses an expensive-to-generate auth mechanism (e.g. a signed URL):
  - server provides an [access\\_id](#)
  - caller passes the [access\\_id](#) to the [/access](#) endpoint
  - server provides an [access\\_url](#) with the generated mechanism (e.g. a signed URL in the [url](#) field)
  - caller fetches the object bytes from the [url](#) (passing auth info from the specified headers, if any)




# Chapter 5. Motivation

Data sharing requires portable data, consistent with the FAIR data principles (findable, accessible, interoperable, reusable). Today's researchers and clinicians are surrounded by potentially useful data, but often need bespoke tools and processes to work with each dataset. Today's data publishers don't have a reliable way to make their data useful to all (and only) the people they choose. And today's data controllers are tasked with implementing standard controls of non-standard mechanisms for data access.



*Figure 1: there's an ocean of data, with many different tools to drink from it, but no guarantee that any tool will work with any subset of the data*

We need a standard way for data producers to make their data available to data consumers, that supports the control needs of the former and the access needs of the latter. And we need it to be interoperable, so anyone who builds access tools and systems can be confident they'll work with all the data out there, and anyone who publishes data can be confident it will work with all the tools out there.



*Figure 2: by defining a standard Data Repository API, and adapting tools to use it, every data publisher can now make their data useful to every data consumer*

We envision a world where:

- there are many many **data consumers**, working in research and in care, who can use the tools of their choice to access any and all data that they have permission to see
- there are many **data access tools** and platforms, supporting discovery, visualization, analysis, and collaboration
- there are many **data repositories**, each with their own policies and characteristics, which can be accessed by a variety of tools
- there are many **data publishing tools** and platforms, supporting a variety of data lifecycles and formats
- there are many many **data producers**, generating data of all types, who can use the tools of their choice to make their data as widely available as is appropriate




Figure 3: a standard Data Repository API enables an ecosystem of data producers and consumers

This spec defines a standard **Data Repository Service (DRS) API** (“the yellow box”), to enable that ecosystem of data producers and consumers. Our goal is that the only thing data consumers need to know about a data repo is *“here’s the DRS endpoint to access it”*, and the only thing data publishers need to know to tap into the world of consumption tools is *“here’s how to tell it where my DRS endpoint lives”*.

## 5.1. Federation

The world’s biomedical data is controlled by groups with very different policies and restrictions on where their data lives and how it can be accessed. A primary purpose of DRS is to support unified access to disparate and distributed data. (As opposed to the alternative centralized model of “let’s just bring all the data into one single data repository”, which would be technically easier but is no more realistic than “let’s just bring all the websites into one single web host”.)

In a DRS-enabled world, tool builders don’t have to worry about where the data their tools operate on lives — they can count on DRS to give them access. And tool users only need to know which DRS server is managing the data they need, and whether they have permission to access it; they don’t have to worry about how to physically get access to, or (worse) make a copy of the data. For example, if I have appropriate permissions, I can run a pooled analysis where I run a single tool across data managed by different DRS servers, potentially in different locations.

# Chapter 6. Paths

## 6.1. Get info about a Data Bundle.

```
GET /bundles/{bundle_id}
```

### 6.1.1. Description

Returns bundle metadata, and a list of ids that can be used to fetch bundle contents.

### 6.1.2. Parameters

Type	Name	Schema
Path	<b>bundle_id</b> <i>required</i>	string

### 6.1.3. Responses

HTTP Code	Description	Schema
200	The Data Bundle was found successfully.	<a href="#">Bundle</a>
400	The request is malformed.	<a href="#">Error</a>
401	The request is unauthorized.	<a href="#">Error</a>
403	The requester is not authorized to perform this action.	<a href="#">Error</a>
404	The requested Data Bundle wasn't found.	<a href="#">Error</a>
500	An unexpected error occurred.	<a href="#">Error</a>

### 6.1.4. Tags

- DataRepositoryService

## 6.2. Get info about a Data Object.

```
GET /objects/{object_id}
```

### 6.2.1. Description

Returns object metadata, and a list of access methods that can be used to fetch object bytes.

### 6.2.2. Parameters

Type	Name	Schema
Path	<b>object_id</b> <i>required</i>	string

### 6.2.3. Responses

HTTP Code	Description	Schema
200	The Data Object was found successfully.	<a href="#">Object</a>
400	The request is malformed.	<a href="#">Error</a>
401	The request is unauthorized.	<a href="#">Error</a>
403	The requester is not authorized to perform this action.	<a href="#">Error</a>
404	The requested Data Object wasn't found	<a href="#">Error</a>
500	An unexpected error occurred.	<a href="#">Error</a>

### 6.2.4. Tags

- DataRepositoryService

## 6.3. Get a URL for fetching bytes.

```
GET /objects/{object_id}/access/{access_id}
```

### 6.3.1. Description

Returns a URL that can be used to fetch the object bytes.

This method only needs to be called when using an [AccessMethod](#) that contains an [access\\_id](#) (e.g., for servers that use signed URLs for fetching object bytes).

### 6.3.2. Parameters

Type	Name	Description	Schema
Path	<b>access_id</b> <i>required</i>	An <a href="#">access_id</a> from the <a href="#">access_methods</a> list of a Data Object	string
Path	<b>object_id</b> <i>required</i>	An <a href="#">id</a> of a Data Object	string

### 6.3.3. Responses

HTTP Code	Description	Schema
200	The access URL was found successfully.	<a href="#">AccessURL</a>

HTTP Code	Description	Schema
400	The request is malformed.	<a href="#">Error</a>
401	The request is unauthorized.	<a href="#">Error</a>
403	The requester is not authorized to perform this action.	<a href="#">Error</a>
404	The requested access URL wasn't found	<a href="#">Error</a>
500	An unexpected error occurred.	<a href="#">Error</a>

#### 6.3.4. Tags

- `DataRepositoryService`

### 6.4. Get information about this implementation.

GET /service-info

#### 6.4.1. Description

May return service version and other information.

#### 6.4.2. Responses

HTTP Code	Description	Schema
200	Service information returned successfully	<a href="#">ServiceInfo</a>

#### 6.4.3. Tags

- `DataRepositoryService`

# Chapter 7. Definitions

## 7.1. AccessMethod

Name	Description	Schema
<b>access_id</b> <i>optional</i>	An arbitrary string to be passed to the <code>/access</code> method to get an <code>AccessURL</code> . This string must be unique per object. Note that at least one of <code>access_url</code> and <code>access_id</code> must be provided.	string
<b>access_url</b> <i>optional</i>	An <code>AccessURL</code> that can be used to fetch the actual object bytes. Note that at least one of <code>access_url</code> and <code>access_id</code> must be provided.	<code>AccessURL</code>
<b>region</b> <i>optional</i>	Name of the region in the cloud service provider that the object belongs to. <b>Example</b> : <code>"us-east-1"</code>	string
<b>type</b> <i>required</i>	Type of the access method.	enum (s3, gs, ftp, gsiftp, globus, htsgget, https, file)

## 7.2. AccessURL

Name	Description	Schema
<b>headers</b> <i>optional</i>	An optional list of headers to include in the HTTP request to <code>url</code> . These headers can be used to provide auth tokens required to fetch the object bytes. <b>Example</b> : <pre>{   "Authorization" : "Basic Z2E0Z2g6ZHJz" }</pre>	< string > array
<b>url</b> <i>required</i>	A fully resolvable URL that can be used to fetch the actual object bytes.	string

## 7.3. Bundle

Name	Description	Schema
<b>aliases</b> <i>optional</i>	A list of strings that can be used to find other metadata about this Data Bundle from external metadata sources. These aliases can be used to represent the Data Bundle's secondary accession numbers or external GUIDs.	< string > array

Name	Description	Schema
<b>checksums</b> <i>required</i>	<p>The checksum of the Data Bundle. At least one checksum must be provided.</p> <p>The Data Bundle checksum is computed over a sorted concatenation of all the checksums (names not included) within the top-level 'contents' of the Bundle (not recursive). The list of Data Object or Bundle checksums are sorted alphabetically (hex-code) before concatenation and a further checksum is performed on the concatenated checksum value.</p> <p>Example below:</p> <p>Data Objects:</p> <p>md5(DO1) = 72794b6d30bc86d92e40a1aa65c880b8</p> <p>md5(DO2) = 5e089d29a18954e68a78ee6a3c6edabd</p> <p>Data Bundle:</p> <p>DB1 = md5( concat( sort( md5(DO1), md5(DO2) ) ) )</p> <p>= md5( concat( sort( 72794b6d30bc86d92e40a1aa65c880b8, 5e089d29a18954e68a78ee6a3c6edabd ) ) )</p> <p>= md5( concat( 5e089d29a18954e68a78ee6a3c6edabd, 72794b6d30bc86d92e40a1aa65c880b8 ) )</p> <p>= md5( 5e089d29a18954e68a78ee6a3c6edabd72794b6d30bc86d92e40a1aa65c880b8 )</p> <p>= f7a29a0422e7d870b10839ad6c985079</p>	< <a href="#">Checksum</a> > array
<b>contents</b> <i>required</i>	The list of Data Objects and Data Bundles contained by this Data Bundle.	< <a href="#">BundleObject</a> > array
<b>created</b> <i>required</i>	Timestamp of Bundle creation in RFC3339.	string (date-time)
<b>description</b> <i>optional</i>	A human readable description of the Data Bundle.	string
<b>id</b> <i>required</i>	An identifier unique to this Data Bundle.	string
<b>name</b> <i>optional</i>	A string that can be used to name a Data Bundle.	string
<b>size</b> <i>required</i>	The cumulative size, in bytes, of all Data Objects and Bundles listed in the <b>contents</b> field.	string (int64)
<b>updated</b> <i>optional</i>	Timestamp of Bundle update in RFC3339, identical to create timestamp in systems that do not support updates.	string (date-time)
<b>version</b> <i>optional</i>	A string representing a version. (Some systems may use checksum, a RFC3339 timestamp, or an incrementing version number.)	string

## 7.4. BundleObject

Name	Description	Schema
<b>drs_uri</b> <i>optional</i>	A list of full DRS identifier URI paths that may be used obtain the Data Object or Data Bundle. These URIs may be external to this DRS instance. <b>Example :</b> <code>"drs://example.com/ga4gh/drs/v1/objects/{object_id}"</code>	< string > array
<b>id</b> <i>required</i>	A DRS identifier of a Data Object or a nested Data Bundle.	string
<b>name</b> <i>required</i>	A name declared by the Bundle author that must be used when materialising the associated data object, overriding any name directly associated with the object itself. This string MUST NOT contain any slashes.	string
<b>type</b> <i>required</i>	The type of content being referenced. BundleObject of type bundle will need to be recursed further.	enum (object, bundle)

## 7.5. Checksum

Name	Description	Schema
<b>checksum</b> <i>required</i>	The hex-string encoded checksum for the data	string
<b>type</b> <i>optional</i>	The digest method used to create the checksum. If left unspecified md5 will be assumed.  possible values: md5 # most blob stores provide a checksum using this etag # multipart uploads to blob stores sha256 sha512	string

## 7.6. Error

An object that can optionally include information about the error.

Name	Description	Schema
<b>msg</b> <i>optional</i>	A detailed error message.	string
<b>status_code</b> <i>optional</i>	The integer representing the HTTP status code (e.g. 200, 404).	integer

## 7.7. Object



Name	Description	Schema
<b>access_methods</b> <i>required</i>	The list of access methods that can be used to fetch the Data Object.	< <a href="#">AccessMethod</a> > array
<b>aliases</b> <i>optional</i>	A list of strings that can be used to find other metadata about this Data Object from external metadata sources. These aliases can be used to represent the Data Object's secondary accession numbers or external GUIDs.	< string > array
<b>checksums</b> <i>required</i>	The checksum of the Data Object. At least one checksum must be provided.	< <a href="#">Checksum</a> > array
<b>created</b> <i>required</i>	Timestamp of object creation in RFC3339.	string (date-time)
<b>description</b> <i>optional</i>	A human readable description of the Data Object.	string
<b>id</b> <i>required</i>	An identifier unique to this Data Object.	string
<b>mime_type</b> <i>optional</i>	A string providing the mime-type of the Data Object. <b>Example</b> : "application/json"	string
<b>name</b> <i>optional</i>	A string that can be used to name a Data Object.	string
<b>size</b> <i>required</i>	The object size in bytes.	integer (int64)
<b>updated</b> <i>optional</i>	Timestamp of Object update in RFC3339, identical to create timestamp in systems that do not support updates.	string (date-time)
<b>version</b> <i>optional</i>	A string representing a version.	string

## 7.8. ServiceInfo

Useful information about the running service.

Name	Description	Schema
<b>contact</b> <i>optional</i>	Maintainer contact info	object
<b>description</b> <i>optional</i>	Service description	string
<b>license</b> <i>optional</i>	License information for the exposed API	object
<b>title</b> <i>optional</i>	Service name	string
<b>version</b> <i>required</i>	Service version	string